From: **Rob Reuss** robreuss@mac.com
Subject: Re: Some CSV files
Date: February 10, 2026 at 6:04 PM
To: Marc Weiss mweiss@weblab.org

I should have been more clear when I sent the CSV files.

The goal was simply to give you a sense of the underlying content, presented in a concise way so you can get a feel of the content we're working with, **with specific attention to user identity**, because it is powerful in a different way than subject-oriented guidance.

'm sorry if it took you down the path of being concerned the original structure was lost. Nothing is lost.

Having the ability to see what other things a person said is a common expectation - it adds a lot of context and understanding to a given post - something as fundamental as biographical details they shared in an earlier post (e.g. "I was a nurse in Saigon…"). It also introduces the possibility of having a view of everything a given person said, with is particularly important in terms of those who were highly engaged.

In general I'm intuitively moving in direction where people will more substantively engage with it, and maybe that's surplus to requirements for you, and you prefer to present it as it was, with maximum fidelity - but that's a mistake for me from a time management standpoint.

Adding more ways of explore the content does diverge from the original, and preservation organizations/sites seem to approach archiving early websites with that sort of attitude (this is what you get with the Wayback Machine) - that is, full fidelity to original site, including design problems and weaknesses in navigation…precisely because those elements are an essential reflection of the early web. How much they do that for principled reasons versus the lack of a budget to provide analytic tools probably varies by archival program. And to what extent they would insist on complete fidelity, which even our more limited plan lacks in a complete sense, I don't know.

So archiving this site is a balancing act between two very legitimate approaches:

1) full fidelity to the original structure and design (because the structure and design are a part of this story of the early web) and

2) providing the kinds of tools for exploring the archive on a semantic level, to enable explorers to develop a more meaningful and deeper understanding of the content itself, because honestly, given the navigation on the original site people will probably read a relatively small number of posts out of curiosity but they are not going to fully explore it.

My app is the absolute priority in my life and I have a ways to go - no fault of yours, I'm letting this project attract too much of my attention (because it is very interesting).

I'm happy to produce a version with maximum fidelity to the original, same basic navigation pattern (the screenshot you sent), with greater portability (as a static HTML site), and get it up on GitHub so you can link to it. I'll minimize use of ChatGPT - my of the script for the generation of the static site is complete.

I'll add the letter as the initial page and user will click through to the site. You should probably reiterate credits in the letter, particularly if you to let original stakeholders know about it - we do not have a plan to put the original credits page to be in the new site. You should probably add a press inquiries e-mail address at the bottom.

Expect something similar to Wayback, but with no broken links and the horribly scaled graphical elements gone, although an LLM service that shall remain nameless was able to produce the attached cleaned and upsized version of the banner over the topics index.

attached cleaned and upsized version of the banner over the topics index.

> On Feb 10, 2026, at 5:23 PM, Marc Weiss <mweiss@weblab.org> wrote:
>
> Found what the dialogue main page looked like on the Wayback Machine and just sent it to you with a screenshot of the first 33 of the 225 topics (and with number of posts and last post displayed after each topic)
>
> So rather than trying to figure out the sequencing in the csv files, can that be recreated?
>
>> On Feb 10, 2026, at 12:22 PM, Rob Reuss <robreuss@mac.com> wrote:
>>
>>> I'll take a look at the new threads file after dinner, but regarding the sort: people were posting on multiple threads over the course of the dialogue, so if won't sorting by post_date result in a hodgepodge of threads?
>>
>> Sorting by post_date will show the posts in the order they were created, so yes, in terms of the sequence of the rows they would be like you said: people post to this thread, then that thread, and back to this thread, interweaved.
>>
>> If you sort on thread_index_id or thread_title, and use post_date as the secondary sort, you would see each thread as a group of rows with the posts within them sorted by the order they were posted.
>>
>> If you prefer a series of files that represent each thread, with the posts within them, I can do that easily.
>>
>> If you come up with a question about the data you want answered just ask - it's easy for me to

do the analysis and/or exports.

> On Feb 9, 2026, at 5:45 PM, Marc Weiss <mweiss@weblab.org> wrote:
>
> It's been a long day on meetings and racing deadlines, so I'm just getting back to your emails now.
>
> Happy to know the index numbers were not identifiers in the original content!
>
> I'll take a look at the new threads file after dinner, but regarding the sort: people were posting on multiple threads over the course of the dialogue, so if won't sorting by post_date result in a hodgepodge of threads?
>
> > On Feb 9, 2026, at 9:47 AM, Rob Reuss <robreuss@mac.com> wrote:
> >
> > The attached is more up to date - includes more content from an additional directory. There are two additional directories that have content that I haven't included yet in the processing - it's unclear from the directory names what they are and I wanted to avoid importing duplicate data.
> >
> > Your perception about which threads are "first" is probably due to sorting the file. If you sort on the post_date then you would be seeing content in the original order of creation.
> >
> > The index numbers in the file were generated as a part of the processing of the files, they are not identifiers in the original content.
> >
> > <threads.csv>
> >
> > > On Feb 8, 2026, at 11:08 PM, Marc Weiss <mweiss@weblab.org> wrote:
> > >
> > > Shared the letter from the producer with you a little while ago. It's long…
> > >
> > > I imported these files to a google spreadsheet and spent a little time going through the threads. I did a sort by thread index, then post id.
> > >
> > > Several things look off. It shows the second thread as Vietnam Protest, but my memory is that there was at least one (maybe more) threads before that. After post 83 in that thread, the thread index goes to 3, but it's the same topic.
> > >
> > > After post 63 in the second Vietnam protest thread, the thread index goes from 3 to 5. Still Vietnam Protest. And where's 4?
> > >
> > > In fact, Vietnam protest goes through line 658 (thread index 7 at that point, but 7 keeps going with ANOTHER topic). Something's not right.
> > >
> > > Maybe I made a mistake sorting them. I'll re-upload the thread index and try again tomorrow.
> > >
> > > Sure hope the file's not corrupted!
> > >
> > > > On Feb 8, 2026, at 6:35 PM, Rob Reuss <robreuss@mac.com> wrote:
> > > >
> > > > Hey Marc,
> > > >
> > > > Here's some CSV files that should open right up in Excel or Numbers, if you're on a Mac:
> > > >
> > > > **threads.csv** is all the structured fields I was able to extract - it is everything:

"post_created_at",
"thread_title",
"post_title",
"post_email",
"post_name",
"post_body",
"post_id",
"thread_index_id"

You should be able to sort on the headers with a left mouse click on the column headers (the letters “E”, “F”, etc., not the field headers.

post_names.csv contains the person name strings that were extracted - there’s **4167** unique strings, but in many cases the same person repeats because the strings are slightly different, because of 1) misspellings, 2) variation in the use of capitalization, and 3) variations between using first and last name, etc.

post_names_dedupped.csv reduces the number of name strings to **1050** by only including one instance for each duplicate name string and doing that comparison as lowercase, so the mismatched capitalizations doesn’t matter. It does not address the other issues - those are things that AI might be able to help with, with human supervision.

1050 is a significant over-estimate of the number of participants - if you sort the list in Numbers you’ll see that there’s still a lot of duplicate persons because of the aforementioned inconsistencies, and some people use slogans as names so it’s hard to say if they posted more than once. But scanning by eye my guess would be that the number of truly unique participants was 500-800, but I could be way off.

Let me know if any additional analysis would be helpful.

Rob

<post_names_dedupped.csv>
<post_names.csv>
<threads.csv>